

Resident OS

Screen reference and demo script

22 screens across 5 surfaces
wireframes, components, data bindings, states

build reference · v1.0 · July 2026
companion to PRD.md · ARCHITECTURE.md · DESIGN.md · PHASES.md

Part one — Screen reference

Every screen in Resident OS, specified enough to build from. Owner: Track B (Surface), with data bindings owned by Track A
Related: DESIGN.md (tokens and language) · PRD.md (requirement IDs) · PHASES.md (when each ships) · DEMO.md (the click path)

How to use this file: each screen has a wireframe, a component list, its **data bindings** (which table or endpoint feeds each element), its states, and acceptance criteria. If you are driving Claude Code, point it at the individual screen section rather than the whole file.

Legend used in wireframes

- filled status lamp ○ hollow lamp (unclassified)
- ◆ decider diamond – machine blue / human brass
- [C1] citation chip ▮ guardrail hit
- ¶ § • authority glyphs – SOP / lease / statute / vendor contract

Screen index

#	Screen	Surface	Phase	Requirement
1	Landing / demo entry	public	5	—
2	Resident home	resident	3	—
3	Report an issue	resident	3	FR-101...107
4	Resident ticket detail	resident	3	FR-701
5	Resident notifications	resident	8	FR-801...806
6	Manager approval queue	manager	3	FR-601...604
7	Trace viewer	manager	5	FR-702, FR-703
8	Retrieval inspector	manager	5	FR-302
9	Council view	manager	6	FR-308...310
10	Context budget	manager	6	—
11	Knowledge browser	manager	2	FR-502
12	Governance tab	manager	4	FR-704, FR-705
13	Agent health	manager	6	—
14	Cost dashboard	manager	5	NFR-06
15	Eval dashboard	manager	5	\$7.1
16	Gateway dashboard	manager	7	—
17	Memory review	manager	6	—
18	Owner monthly dashboard	owner	4	FR-705
19	Tech job card	tech	3	FR-401
20	Vendor accept/decline	vendor	3	FR-404, FR-405
21	Compliance block	manager	4	FR-504, FR-508
22	Degraded mode banner	all	3	NFR-10

1 · Landing / demo entry

Job: get a stranger into a live demo in under 15 seconds, with no form.



Hero rule: the hero is the product doing its one trick — an auto-playing (muted, reduced-motion-safe) trace replay. Not a big number with a gradient.


Data: GET /v1/runs/{seed_run_id} replayed client-side from a recorded event fixture, so the landing page never depends on a live model call.

States: reduced-motion → static completed trace. JS disabled → static screenshot.

Acceptance: a first-time visitor reaches a working manager console in ≤ 2 clicks with no signup.

2 · Resident home

Job: answer "is my thing being handled" in one glance, then get out of the way.



Why asymmetric bento: tile size encodes priority of attention. A uniform grid would say the community events matter as much as the open plumbing ticket. They do not.

Data bindings

Element	Source
Greeting, unit	people, tenancies
Active ticket tile	GET /v1/tickets?status=open&limit=1
Events strip	events (Phase 9; hide the strip entirely before then)
Marketplace / amenities	Phase 8 — render as disabled tiles before then

States: no open ticket → the report tile expands to full width. New resident, no history → same, plus a one-line welcome.

3 · Report an issue

Job: submitted in under 30 seconds, three taps, no dropdowns.

STEP 1 What's wrong?

Water	Heat &	Power
Appliance	Pests	Something else

large tiles, plain language

[Tell us more]

STEP 2 Show us

[camera-first sheet · up to 5 photos]

[● hold to record · 30s max]

STEP 3 Access

Can we enter if you're out? (Yes) (No)

Pets in the apartment? [toggle]

Best time [morning][afternoon][evening]

[Send]

Non-negotiable (FR-102): the confirmation with a ticket number appears in under 1 second and **never waits on a model call**. Enqueue, then acknowledge.

Data bindings

Element	Source
Submit	POST /v1/tickets with Idempotency-Key header
Photos	signed R2 upload URL, then ticket_media
Voice	upload → transcribe in worker → raw_text appended, audio deleted

States: offline → queued locally with a visible "will send when you're back online" chip. Upload failure → ticket still submits, media retries.

Acceptance: double-tapping Send creates exactly one ticket.

4 · Resident ticket detail

Job: show what is happening in plain language, and explain money before it is asked about.

< #4471

● P2 · URGENT

Kitchen sink

Reported 23 Jul, 11:40pm

✓ We received your report11:40pm

✓ Checked against your lease11:40pm

✓ Marked urgent – third time11:41pm

this year

● Plumber scheduled7:15am

Delta Plumbing · before 10am

○ Repair

There's no charge to you for this repair.

[Why? ▼]

↳ Your lease (section 7.3) puts normal wear and tear on the owner. This drain has been repaired twice before, so it counts as wear, not damage.

[Add a photo]

[Message us]

Design rules - The resident sees **no** confidence score, **no** model name, **no** agent names. - Same underlying events as the engineering trace, different vocabulary — "Checked against your lease," not "policy retrieval." - "Why?" is collapsed by default and expands to the plain-English reason **with the lease section named**. A resident who can see why they are not being charged trusts the system.

Data bindings

Element	Source
Timeline	agent_steps mapped through a resident-facing label map, not rendered raw
Charge explanation	decisions.final.responsible_party + the citations entry whose authority = lease
Vendor + window	work_orders

States: decision still running → timeline shows completed steps with the next one pulsing. Escalated to a human → "A team member is reviewing this" with no fake progress.

5 · Resident notifications (Phase 8)

Notifications	[⦿]
● Plumber arriving before 10am #4471 · kitchen sink	7:15am
Rent due in 3 days	Mon
3 community updates	Sun

← batched U3

SETTINGS
Quiet hours [10pm]-[7am]
Urgent overrides quiet hours [on, locked]
Community updates [weekly digest ▼]

Rule (FR-802): the settings screen governs a **deterministic policy engine**. The U0 override toggle is visibly locked on — life-safety alerts are not a preference.

Data: `notifications`, `notification_budgets`. Suppressed items are visible to staff in the governance tab but not to residents.

6 · Manager approval queue — the hero screen

Job: clear the queue without making an expensive mistake. Marcus is here four hours a day.

QUEUE 3 awaiting approval

[SLA risk ▼] [All properties]

● #4471 4B · Kitchen plumbing Third drain ticket in 14 months. SOP-PLM-04 escalates to trap replacement after two repeats. Lease 7.3: normal wear, owner pays. Delta Plumbing · est. \$180 · owner-billed [Approve] [Edit] [Reject]	SLA 13h conf 0.91	◆
● #4468 2C · No hot water Council: members disagreed on priority	SLA 4h conf 0.74	◆
○ #4463 1A · Cabinet door Escalated — the system declined to decide	SLA 4d conf 0.44	◆

REJECT expands → Wrong priority | Wrong trade | Wrong party | Policy misread | Other — the eval label taxonomy

Five decisions worth defending in an interview

- Sorted by SLA risk, never arrival time.** The queue exists to prevent the expensive miss.
- One card, one screen, no scroll.** If the manager scrolls to decide, the card is too big.
- The diamond changes hands.** Blue while the system owns it; brass the instant a person touches it. Consistent across every surface.
- Council disagreement is on the card, not three clicks deep.** Disagreement is the highest-information thing the system can tell a manager.
- Every button is a labelled training example.** Reject → five-option taxonomy → `decisions.override_reason`.

Keyboard map: `j/k` move · `a` approve · `e` edit · `r` reject then `1-5` for reason · `t` trace · `/` search · `?` shortcuts.


Data bindings

Element	Source
Card list	GET /v1/approvals?assignee=me ordered by <code>tickets.sla_due_at</code>
Lamp + priority	<code>decisions.proposed.priority</code>
Reasoning line	<code>decisions.proposed.summary</code> — never the raw model output
Confidence	<code>decisions.confidence</code> (calibrated, not self-reported)
Diamond	<code>decisions.mode</code>
Council flag	<code>agent_runs.council_used</code> + a reviewer dissent flag
Approve	POST /v1/approvals/{id}/approve → issues a single-use <code>approvals.token</code>

States: empty → "Queue clear. 3 tickets resolved today." Never a shrug illustration. Optimistic approve with rollback on failure.

7 · Trace viewer

Job: answer "why did it do that" for an engineer, a manager, or an auditor.




Build notes - Virtualise the tree. Agent traces nest deeply; a tree that stutters on expand destroys debugging speed, which is the only reason this screen exists. - **Parallel spans share a horizontal track** so "three members ran at once" is visible rather than asserted. - Steps with no model call (context.broker , context.budget , decision.gate) are labelled code . This makes the ≤ 5 -LLM-call budget visible, which is the single best thing to point at in an interview. - Replay → GET /v1/runs/{id}/replay re-executes against pinned prompt, model, and KB versions and shows a side-by-side diff.

Data: agent_runs , agent_steps , llm_calls , retrievals , guardrail_events — one query, assembled server-side into a nested tree.

8 · Retrieval inspector

Job: answer "why was the wrong chunk in the window."




The most useful element is the budget cut line. Chunks below it are dimmed but present — what was *almost* retrieved is usually the answer to "why was this wrong."

Data: retrievals.results — store bm25_rank , vec_rank , rrf , rerank_score , and used per chunk at write time. If you do not store these in Phase 2, this screen cannot be built in Phase 5.

9 · Council view

Job: show that disagreement carries information.




The Blind to line is the whole argument. Council members differ by evidence, not persona, because same-model samples over identical context produce correlated errors — they agree confidently and wrongly. Make the blindness visible or the design reads as three personas.

Style rule: when members disagree, the disagreement is the **headline** of the screen, not a warning banner.

Data: agent_steps where agent LIKE 'council.member.%', plus the synthesis step and the deterministic reviewer's flags.

10 · Context budget



Hatching for reserved space is a drawing convention and reads instantly. Slots chronically at 40% utilisation are over-provisioned — reallocate them.

Data: envelope metadata stored on the `context.budget` step's `output`.

11 · Knowledge browser




Showing precedence on the screen matters: it tells a manager that a conflict between the SOP and the statute is resolved by rule, not by the model's judgment.

Data: `kb_documents`, `kb_chunks`. Search hits `GET /v1/knowledge/search` and returns scores, which doubles as the demo surface.

12 · Governance tab

Job: the screen that makes an institutional buyer take a second meeting.




The **override-reason distribution** is a product feature, not an admin log. It tells an operator where the system is weak, and it tells you where to spend the next eval cycle.

Export PDF is designed for print — light theme, hairline rules, page furniture, title-block header on every page. This is the artifact an operator's counsel attaches to a file.

Data: `decisions`, `guardrail_events`, hash-chain verification job output.

13 · Agent health




Small multiples, uniform card size, so anomalies pop by **shape** rather than by scale.

Alert-worthy, not just displayed: escalation rate moving $\pm 50\%$ week over week in *either* direction. A drop means the system got overconfident, which is worse than a rise.

Data: aggregates over `agent_steps` and `llm_calls`, rolled up nightly by the worker.

14 · Cost dashboard




The ceiling is drawn as a **rule**, not signalled by a colour change — consistent with the drawing language, and it survives colour-blind viewing.

Data: `llm_calls` aggregated; cache hit rates from the gateway policy layer.

15 · Eval dashboard


Job: the screen that makes a hiring manager believe everything else.



Design rules - The missed metric is styled as missed and stays on the dashboard. Publishing a red number with an explanation is more credible than five green ones, and it is the fastest signal that you have operated something. - The **calibration reliability diagram is the hero chart** — give it the largest tile. Almost no portfolio project has one. - Always show the git SHA, the pinned judge model, and the seed. An unreproducible metric is not a metric.

Data: `eval_runs.metrics`, `eval_items`, calibration bins computed at eval time and stored.

16 · Gateway dashboard (Phase 7)



Judge deliberately routes to a different model family than generation, to avoid self-preference bias. Show that on the screen — it is a question you will be asked.

17 · Memory review (Phase 6)

PROPOSED REFLECTIONS

"Riverside building C has low water pressure, causing recurring cartridge failures in single-handle faucets." scope: building C · confidence 0.81 evidence #3102 #3877 #4102 #4471 [Approve → promote] [Reject] [Needs more evidence]

Approve is a brass button. Memory only becomes durable when a person says so — that is the anti-poisoning design, made visible.

Data: reflections where status = 'proposed'. Approval sets status='approved' and approved_by.

18 · Owner monthly dashboard

RIVERSIDE · March 2026 · 412 units

Cost per unit SLA met Preventive Reopen rate
\$54 94.2% 31% 6.1%
down 8% up 2.1pt up 4pt down 1.2pt

RECURRING FAILURES · CAPITAL REPLACEMENT CANDIDATES
4B Kitchen drain 3 events / 14mo \$540
2C Water heater 2 events / 8mo \$310
7A HVAC blower 4 events / 18mo \$880

AI GOVERNANCE

412 decisions

auto 46% | approved 32% | over 14% | 8%

✓ 0 fair-housing blocks reached a resident

✓ 412 of 412 decisions carry a complete evidence chain

✓ audit log hash chain verified nightly

[Export audit bundle] [Override reasons]

Light theme, print-adjacent, monthly cadence — optimised for comprehension, not interaction. The governance strip gets its own bordered region rather than a tab, because it is the block no competitor ships.

19 · Tech job card

● P2 #4471

RIVERSIDE · 4B

Kitchen sink · plumbing

LIKELY

P-trap failure, third event. Replace assembly — do not snake.

BRING

1½" P-trap kit

Slip-joint pliers

Bucket + towels

ACCESS

Entry permitted · cat in the apartment

[ON MY WAY]

[CAN'T MAKE IT]

The BRING block gets equal weight with the diagnosis. First-time-fix rate is the metric this screen exists to move, and parts-on-truck is the biggest lever on it.

56px minimum touch targets, works with gloves, cached job list with queued status updates when offline.

20 · Vendor accept / decline

Delta Plumbing

Kitchen drain · trap replacement

Riverside, 1400 River Rd, Unit 4B

Before 10:00am, Fri 24 Jul

Access: entry permitted, cat inside

Likely parts: 1½" P-trap assembly

[photo] [photo]

[Accept]

[Decline]

DECLINE →

Not available

Outside my area

Wrong trade

Rate too low

Other

No account, no login, no app. HMAC-signed URL, 48-hour TTL, single-use nonce, no PII in the URL. Decline reasons feed `vendor_scores`, which feeds the next dispatch decision. Closed loop.

21 · Compliance block — the closing argument

Job: demonstrate judgment, not capability. Fifteen seconds of screen time.

▸ Response blocked before sending

Resident asked:
"I use a housing voucher — is that a problem for this unit?"

Guardrail `fair_housing · pre_send`
Reason `source-of-income is a protected class in this jurisdiction; automated response not permitted`
Action `routed to a human · logged · resident told a team member will reply`

Nothing was sent. Nothing was rewritten.

CI · `red-team parity suite` 42 of 42 passing
voucher holder · wheelchair access · service animal · family with children · non-English name

This screen should not look impressive. That is the point. The system did nothing, loudly, and logged it. Pair it on screen with the CI job showing paired protected-class probes passing — that is exactly the test a fair-housing testing organisation would run against you, so you run it against yourself every commit.

22 · Degraded mode banner

▸ Running in reduced mode — decisions are rule-based and all tickets go to review. Tickets are still being accepted normally.

Persistent, non-dismissible, present on every surface including the resident app (in resident vocabulary). Honesty about degradation is a trust asset, and the banner is proof that the degradation ladder exists rather than being claimed.

Cross-cutting requirements

Every screen ships with all five states designed: empty, loading, error, degraded, and blocked-by-guardrail. Skipping these is how a demo falls apart in front of a stranger.

State	Rule
Empty	An invitation with exactly one action. Never a shrug illustration
Loading	Skeleton matching final layout. For runs, show the step structure immediately with pending steps hollow — structure is information before content arrives
Error	What happened + what to do. <i>"Couldn't reach the model provider. The ticket is saved and queued — we'll retry automatically."</i> Errors do not apologise
Degraded	Screen 22
Blocked	Screen 21 — never silently rewritten

Accessibility, checked in CI with axe: WCAG 2.2 AA on both themes · every lamp paired with a text label so colour never carries meaning alone · full keyboard traversal of the queue with a visible focus ring · semantic landmarks · live regions for streaming steps, polite and throttled · `prefers-reduced-motion` honoured · 360px minimum viewport · 44px touch targets (56px on tech).

Data honesty rule: every dashboard renders real seeded data. No placeholder numbers ship, not even temporarily. Experienced reviewers can tell, and it costs all your credibility at once.

Part two — Demo script

How to demo Resident OS, what has to be real, and what to check before anyone sees it. Owner: both tracks · Last updated: 2026-07-25 **Related:** [PRD.md §9](#) (release criteria) · [DESIGN.md](#) (surfaces) · [PHASES.md](#) (Phase 5 is the demo-ready gate)

The demo is a deliverable with its own definition of done, not a thing you improvise on the day. Phase 5 is not complete until every checklist in §3 passes on the production URL.

1. What the demo has to prove

Different audiences need different things proven. The click path is mostly the same; the narration is not.

Audience	What they're deciding	The one moment that decides it
Recruiter / hiring manager	Is this person worth an interview	The eval report with a red number on it
Technical interviewer	Do they understand what they built	Council disagreement + the abstention card
Investor	Is there a company here	The compliance flip and the governance strip
Operator / builder	Would this save me money	The queue, with zero AI vocabulary

Everything below is built around those four moments.

2. Demo fixtures

All demo content is seeded synthetic data, generated by [infra/seed/](#). Never real resident data ([RULES.md §5](#)).

Property: Riverside · 412 units · 6 buildings · jurisdiction [US-VA](#) · 1,200 historical tickets · 40 assets · 8 vendors · 12 knowledge documents.

2.1 The five demo tickets

Each exists to demonstrate exactly one thing. Memorise the IDs.

ID	Unit	Scenario	Demonstrates	Expected outcome
#4471	4B	Water under kitchen sink, smells bad, photo attached	The happy path: retrieval, recurrence detection, citation-grounded decision	P2 · owner-billed · trap replacement · confidence 0.91 · human approval
#4468	2C	No hot water	Council disagreement on priority, resolved toward the more urgent	P1 · warranty claim first · confidence 0.74 · dissent recorded
#4463	1A	Cabinet door won't close	Abstention. Ambiguous lease clause, no unit history	Escalated · confidence 0.44 · responsible party undetermined
#4480	3F	"I smell gas in the hallway"	P0 path with no model between detection and paging	P0 · on-call paged in under 60s · quiet hours ignored
#4455	—	"I use a housing voucher — is that a problem for this unit?"	The compliance flip. Fair-housing guardrail blocks and routes to a human	Blocked · logged · no generated response

2.2 Pre-filled prompt buttons

The landing page carries five buttons that submit these verbatim. Nobody types during a demo.

"There's water under my kitchen sink and it smells bad"
"No hot water since this morning"
"Cabinet door in the kitchen won't close properly"
"I smell gas in the hallway near 3F"
"I use a housing voucher – is that a problem for this unit?"

2.3 Demo accounts

Three buttons on the landing page, pre-authenticated, no form, no signup:

Resident → [priya@demo.residentos.app](#) · unit 4B
Manager → [marcus@demo.residentos.app](#) · Riverside
Owner → [dana@demo.residentos.app](#) · read-only

Vendor view is reached by the signed link on ticket #4471 — demonstrate it, don't make anyone log in for it.

3. Pre-flight checklist

Run this before every public link, every recorded demo, and every interview.

3.1 Hard blockers — do not demo if any fails

- [] `DEMO_MODE=true` in production env, verified by triggering a notification and confirming zero outbound delivery
- [] No real phone number, email, or address exists anywhere in seeded data
- [] Supabase is on Pro (the free tier pauses after 7 idle days — a paused database mid-demo is unrecoverable)
- [] All five demo tickets resolve to their expected outcomes on a fresh run
- [] The eval report at `/evals` matches the committed `docs/evals/REPORT.md` and both cite the same git SHA
- [] Red-team fair-housing suite is green in the latest CI run
- [] No placeholder or hardcoded number appears on any dashboard

3.2 Soft checks

- [] p95 decision latency under 20s measured in the last hour (cold Fly machines add several seconds — hit the URL once, five minutes before)
- [] Model provider status pages clear
- [] Backup demo video accessible offline
- [] Both light and dark themes render correctly on the machine you're presenting from
- [] Browser zoom at 100%, notifications silenced, only the demo tab open

3.3 Reset

```
make demo-reset # truncates runs/decisions/tickets, re-seeds, re-embeds KB, warms caches
```

Run it before every session. A queue with yesterday's approvals already actioned is a flat demo.

4. The 2-minute recruiter demo

The default. Assume they will not click anything themselves and will leave after 30 seconds if nothing happens.

Time	Screen	Action	What you say
0:00	README	—	<i>Nothing.</i> Let the streaming-trace GIF play. Silence here is correct.
0:10	Landing	Click Resident	"This is a resident at 11:40pm with a leak."
0:15	Resident submit	Click prompt 1, submit	"Acknowledged instantly — that never waits on a model. The reasoning runs behind it."
0:30	Resident ticket #4471	Expand "Why?"	"It's telling her <i>why</i> she isn't being charged, and naming the lease clause."
0:45	Landing → Manager	Open queue	"Sorted by deadline risk, not arrival time."
0:55	Queue	Point at card 3 (#4463)	"Look at the third one. Confidence 0.44 — it declined to decide and said exactly what it was missing."
1:10	Trace on #4471	Open trace	"Every step, prompt hash, cached tokens, cost. Three of these have no model call at all."
1:25	Council on #4468	Open council view	"Three specialists with deliberately different evidence. They disagreed, and the disagreement was the useful part."
1:35	Resident → prompt 5	Submit voucher question	"Now watch what happens here."
1:45	Guardrail block	Show block + CI job	"It refuses, logs it, routes to a human. And that's the CI job running protected-class probes on every commit."
1:55	<code>/evals</code>	—	"Every number, including the one I missed."

Stop at 2:00. Do not keep going. If they're interested they will ask, and answering a question is worth more than the next screen you would have shown.

5. The 5-minute technical interview demo

Same path, different emphasis, and you talk about constraints rather than features.

Open with the constraint, not the product:

"Reliability compounds — ten LLM steps at 95% each is 60% end to end. So I budgeted five model calls on the critical path and made everything else deterministic. Three in the normal path, five when council fires."

Then:

1. **Trace viewer (90s).** Walk the span tree. Point out `context.broker`, `context.budget`, and `decision.gate` are code, not model calls. Show the citation verification block: "that's deterministic — I don't need a judge to tell me a citation doesn't resolve."
2. **Retrieval inspector (60s).** BM25 list, vector list, RRF merge, rerank delta, the token-budget cut line. "Embeddings smear identifiers — SOP-PLM-04 ranked 47th in dense-only. Hybrid took recall@5 from 0.62 to 0.84."
3. **Council (90s).** The "blind to" line under each member. "Three samples of the same model over the same context aren't independent — they agree confidently and wrongly. Splitting the evidence makes disagreement informative." Then: "and I measure whether it earns its place. Lift is +0.043 on priority F1. If it were flat, I'd delete it."
4. **Eval dashboard (60s).** Calibration curve, tolerance bands, pinned judge. "The judge is a regression detector, not ground truth — published agreement with human raters is 85 to 92%. The labelled set is the anchor."
5. **Degradation (30s).** "Let me kill the API key in staging." Show rules-only mode. "The building keeps running when a provider doesn't."

Have ready when they push back: why not CrewAI, why Postgres over a vector DB, why you didn't buy a memory vendor, how injection is handled architecturally rather than by scanning, and what the weakest part is. That last answer matters most — see `04-execution-and-career.md` §5 Q12.

6. The 3-minute investor demo

No architecture. Numbers, then the wedge.

1. **Problem, 30s.** "3.3 maintenance requests per unit per year, about \$200 each. Triage is a human guessing at urgency, trade, and who pays. Wrong one way is a \$150 truck roll; wrong the other way is a habitability statute."
2. **Product, 60s.** Queue → one decision → its citations. "Every decision comes with its receipts."
3. **The wedge, 45s.** The compliance flip. "HUD put AI inside the Fair Housing Act in 2024. Testing organisations can probe a live leasing bot remotely, at scale, in an afternoon. Every vendor says they're compliant — what they mean is a line in a system prompt. This is a CI job with protected-class probes."
4. **Governance strip, 30s.** Owner dashboard. "412 decisions, complete evidence chain on every one, exportable. Nobody else ships this."
5. **Economics, 15s.** "Inference is about \$73 per property per year against a \$12,000 account. Margin isn't the constraint here — distribution is."

Expect and answer honestly: "What's your moat?" → "Right now, nothing durable. The compounding assets are the labelled decision corpus and the jurisdiction-mapped guardrail library, and neither exists until there are customers." Naming your non-moats reads as credible; claiming five reads as naive.

7. The operator / builder demo

Rule: do not say "agent," "LLM," "RAG," "council," or "confidence score." Not once.

1. Manager queue. "Eleven decisions waiting, ordered by what's closest to breaching."
2. One card. "Third drain call in fourteen months. It caught the pattern and recommended replacing the trap instead of snaking it again — that's the difference between a \$90 repeat visit and fixing it."
3. Approve. "Five seconds."
4. Card 3. "This one it wasn't sure about, so it asked you. That's deliberate."
5. Owner dashboard. "Cost per unit, SLA compliance, and the units that keep failing — with what you've spent on each."
6. Vendor link. "No app to install. Two taps."

Then stop and ask: "What's the part of your day this doesn't touch?" The answer is worth more than the rest of the demo.

8. The compliance flip, in detail

The most important fifteen seconds, and the easiest to get wrong.

What must happen: the system produces **no answer**. It shows a guardrail block naming `fair_housing`, writes a `guardrail_events` row, routes to a human queue, and the resident sees a neutral "a team member will follow up shortly."

What must never happen: a generated response that discusses source of income, even a careful one. If your system produces a well-worded reply here, the guardrail failed and you should not demo it until it doesn't.

Show alongside it: the CI run of `evals/datasets/redteam_fh.jsonl` — paired probes identical except for a protected-class signal, asserting response parity on latency class, information completeness, and tone.

Say: "This is exactly the test a fair-housing testing organisation would run against a live bot. I run it against myself on every commit."

Do not turn this into a lecture on housing law. Fifteen seconds, then move.

9. Real vs seeded

State this on the landing page. Rigour reads as rigour; hiding it reads as fabrication when someone notices.

Element	Status	Note
Agent runs, traces, prompt hashes, token counts, costs	Real	Live execution against real model APIs
Retrieval scores, RRF ranks, rerank deltas	Real	Computed at query time
Eval metrics, calibration curve, council lift	Real	From a committed run with a git SHA
Guardrail verdicts	Real	Live
Abstention on #4463	Real	Genuine threshold behaviour, not a hardcoded card
Property, units, residents, tenancies	Seeded synthetic	Generated, no real people
1,200 historical tickets, vendor scorecards	Seeded synthetic	Realistic distributions
Owner dashboard month-over-month figures	Seeded synthetic	Derived from seeded history

Landing page line: *"Synthetic property — 412 units, 1,200 historical tickets. All AI execution, retrieval, and evaluation is live."*

If you fake one thing, fake nothing. The abstention card is the best moment in the demo precisely because it's real behaviour; hardcoding it would make every other claim suspect the moment someone probed it.

10. When it breaks live

It will, eventually. Recovery is part of the demo, and handling it well is a better signal than a clean run.

Failure	Say this	Then
Provider slow or down	"This is the degradation path — it drops to rules-only and everything routes to review."	Show it. This is a <i>better</i> demo moment than the happy path.
Cold start delay	"Fly machines scale to zero. First request wakes them."	Keep talking; don't refresh.
Decision differs from expected	"Non-deterministic — let me show you the trace and we can see why."	Open the trace. Never pretend it went as planned.
Total outage	"Let me show you the recording, and we can walk the code instead."	Backup video, then the repo.

Never apologise repeatedly, never blame the network, never say "it usually works." Show the failure path you designed for. That is the point of having designed one.

11. Demo mode specification

`DEMO_MODE=true` must:

- Block every outbound channel — push, email, SMS, voice — at the transport layer, not by config the app reads. A notification renders in the UI and is written to `notifications` with `status='demo_suppressed'`.
- Pin temperature and seeds so runs are reproducible across sessions.
- Disable all destructive admin endpoints.
- Show a persistent, non-dismissible strip: *"Demo environment — synthetic data, no messages are sent."*
- Log every run with `demo=true` so demo traffic never pollutes eval or metrics baselines.

CI asserts that with `DEMO_MODE=true`, no transport adapter is reachable. This is a test, not a convention.

12. Recordings

Two artifacts, both produced in Phase 9.

README GIF — 90 seconds, silent, looping, under 8MB. Submit → stream → decision with citations → approve → the diamond changing from machine to human. No cursor wandering, no dead air. This is what 80% of visitors will see and the only thing many of them will see.

Loom — 3 minutes, narrated. The 2-minute script plus 60 seconds of architecture at the end. Linked from the README, from every application, and from LinkedIn.

Do not record a 20-minute walkthrough. Nobody watches it, and its existence signals you don't know what matters.

13. Anti-patterns

- Making them sign up, fill a form, or read before they can see anything
- Narrating the architecture during the recruiter demo — show the product, save architecture for when they ask
- Autoplaying audio
- Opening with the trace viewer (it's the payoff, not the hook)
- Demoing on a machine with notifications on
- Saying "this is just a prototype" — the eval report is the appropriate caveat and it's a stronger one
- Claiming a customer, a pilot, or an outcome that doesn't exist
- Going past your time budget because it seems to be going well